



Xilinx Technology Showcase & Hackathon 2017

Connect

•

Learn

•

Share

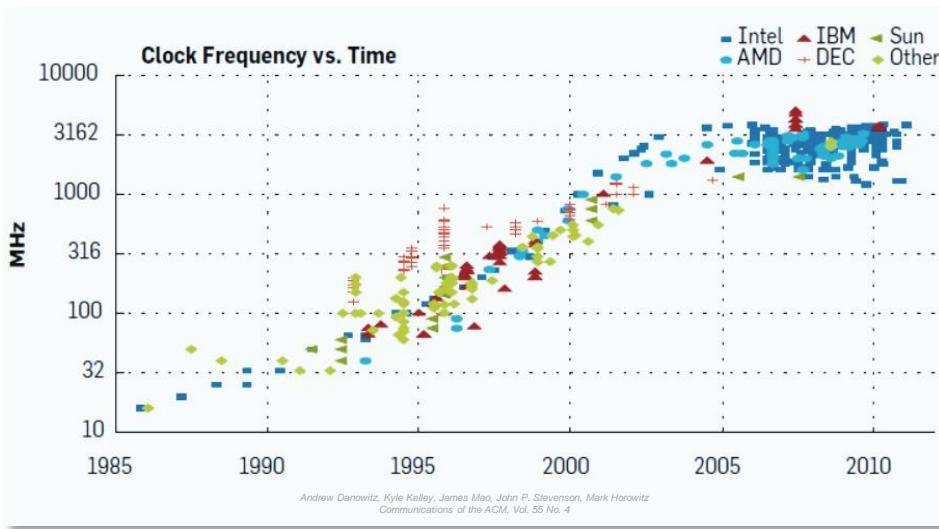


Welcome to the
Xilinx Hackathon 2017

CPU Architectures not Scaling with Workloads

Continued scaling requires accelerated computing, from Edge to Cloud

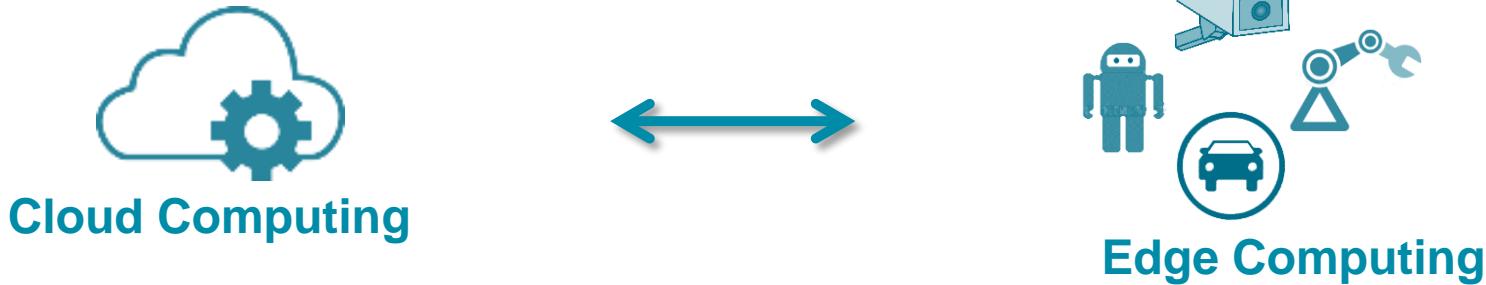
- Processor frequency scaling ended in 2007
- Multicore architecture scaling has flattened



- Workloads require higher performance, lower latency
 - Cloud: video, big data, AI...
 - Edge: auto, surveillance, AI...

Workloads Move from Edge to Cloud... and Back

Reconfigurable accelerators also needed at the edge



Applications are moving to the Cloud:
centralized, elastic resources...

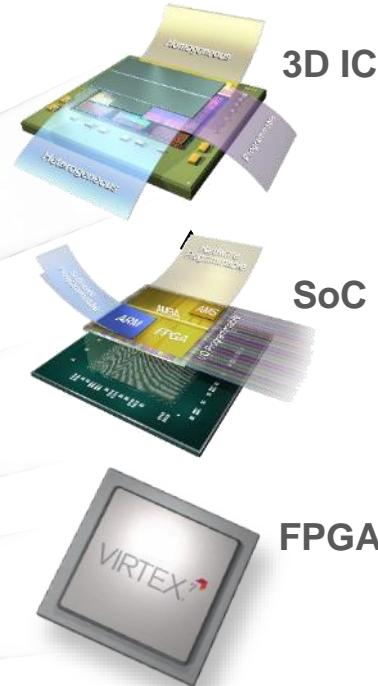
Applications are moving to the Edge:
preprocessing, lower network transfer

Charting an Aggressive Course Forward

Programmable Systems Integration

Break Out

Surpassing Moore's Law
Unprecedented Capacity & Performance
System Intelligence & Integration
Silicon + Architecture for Price/Performance/Watt



Tools

VIVADO[®]

UltraFAST[™]
Design Methodology

SDx[™]
Environments

Architecture

7 Series

UltraScale[™] Architecture

Process

28HPL

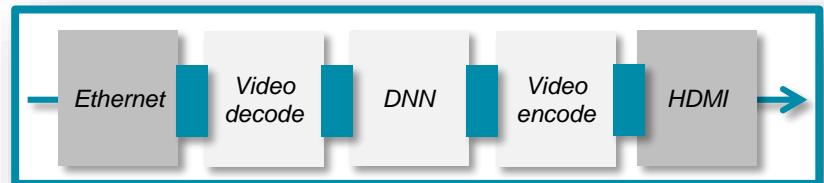
20SoC

16FinFET+

7nm

Emerging Solution: All Programmable Accelerators

- Custom, accelerated processing
 - Reconfigurable to handle diverse workloads
- Custom memory hierarchy
 - Avoid excessive DDR transfers
- Custom IOs
 - Sensor Fusion at the edge



Custom dataflow, memory, IO



Virtex®
FPGA



Zynq®
MPSoC

Harness the Flexibility of Programmable HW

➤ HW engineers



➤ Embedded SW developers



➤ Application SW developers



OpenCL



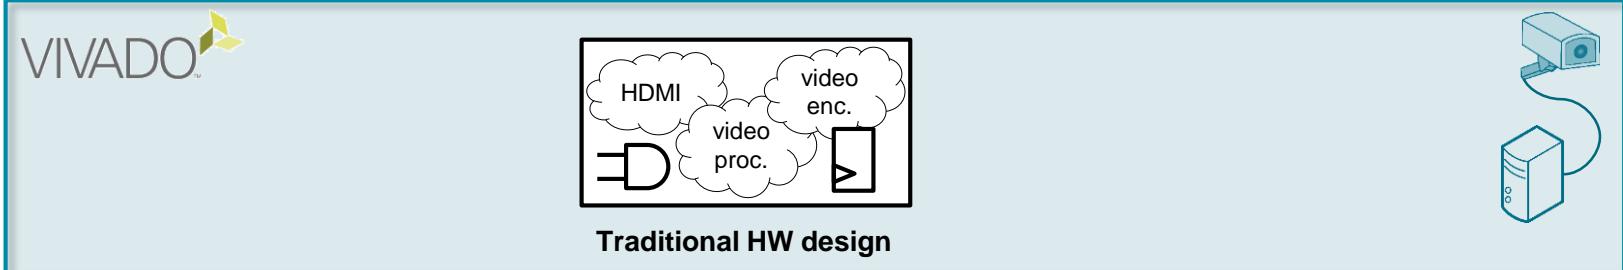
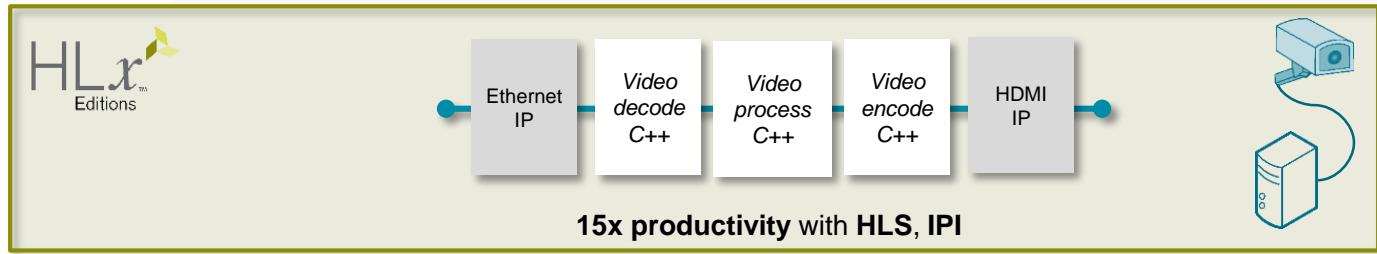
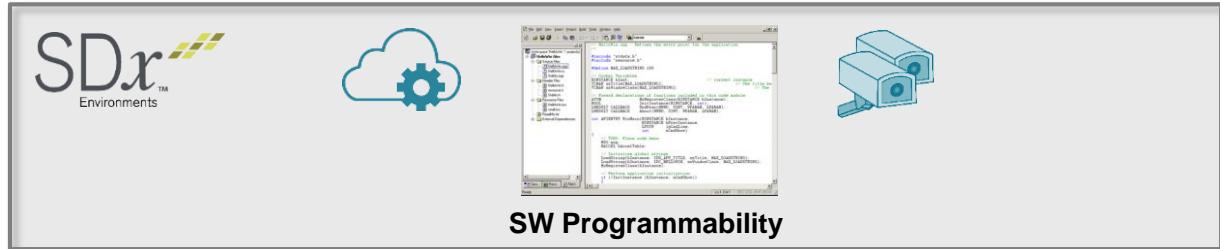
Virtex®
FPGA



Zynq®
MPSoC

Breakout in Programming Model

Development Productivity



Xilinx Product Families: A Broad Portfolio

SoC Portfolio

Cost-Optimized → Mid-Range



All Programmable SoC

System Optimized for Power & Cost

Mid-Range → High-End



All Programmable Heterogeneous MPSoC

Right Engines for the Right Tasks

FPGA & 3D IC Portfolio

Cost-Optimized



Lowest Power & Cost

Mid-Range



Price/Performance/Watt

High-End



Performance & Capacity

All Programmable SoC - Zynq®-7000

Zynq-7000S

Single-Core @ 766MHz
Artix-7 FPGA Fabric



Application Processor

- 32-bit ARMv7
- Up to 1GHz

Zynq-7000

Dual-Core @ 800MHz
Artix-7 FPGA Fabric

Zynq-7000

Dual-Core @ 1GHz
Kintex-7 FPGA Fabric



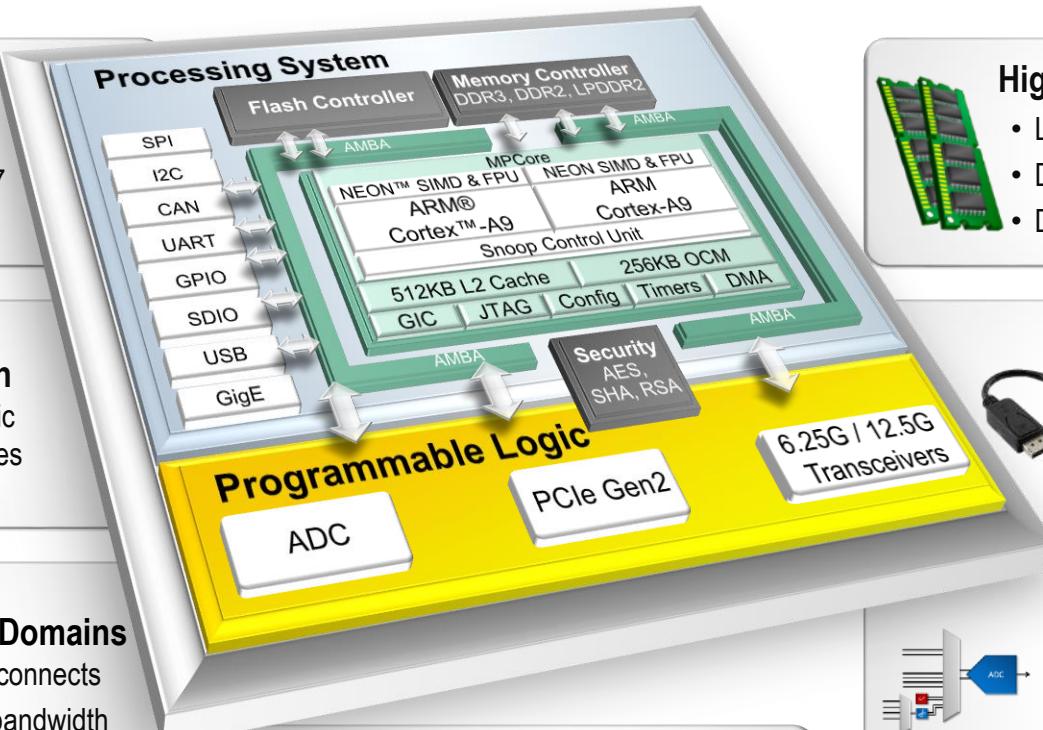
FPGA Acceleration

- 7 Series FPGA Fabric
- Customizable Engines



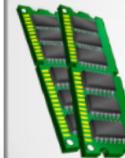
Tightly Coupled Domains

- 3000+ PS/PL interconnects
- Up to 100Gb/s of bandwidth



High Bandwidth Memory

- L1/L2 CPU Caches
- Dedicated On-Chip Memory (OCM)
- DDR3, DDR2, LPDDR2 w/ ECC



Integrated Memory Mapped Peripherals

- USB2.0, GigE, CAN
- UART, SDIO, I2C, SPI



Integrated Analog

- Dual multi-channel 12-bit ADC
- Temp & Voltage sensors



Extensive IP Portfolio

- Standardized AXI4 interfaces
- Enables peripheral expansion



Welcome and Orientation

- Logistics
- Hackathon Theme & Goals
- Community Projects
- Head back up to Retreat!



Organizers



Craig Abramson



Mindy Brooks



Derek Curd



Greg Daughtry



Brad Fross



Adrian Hernandez



Graham Schelle

Helpers

► We will be with you the full 30 hours (in shifts)!



Kv Thanjavur Bhaaskar



Anurag Dubey



Vikhyat Goyal



David Gronlund



Dustin Richmond



Naveen Purushotham



Rock Qu

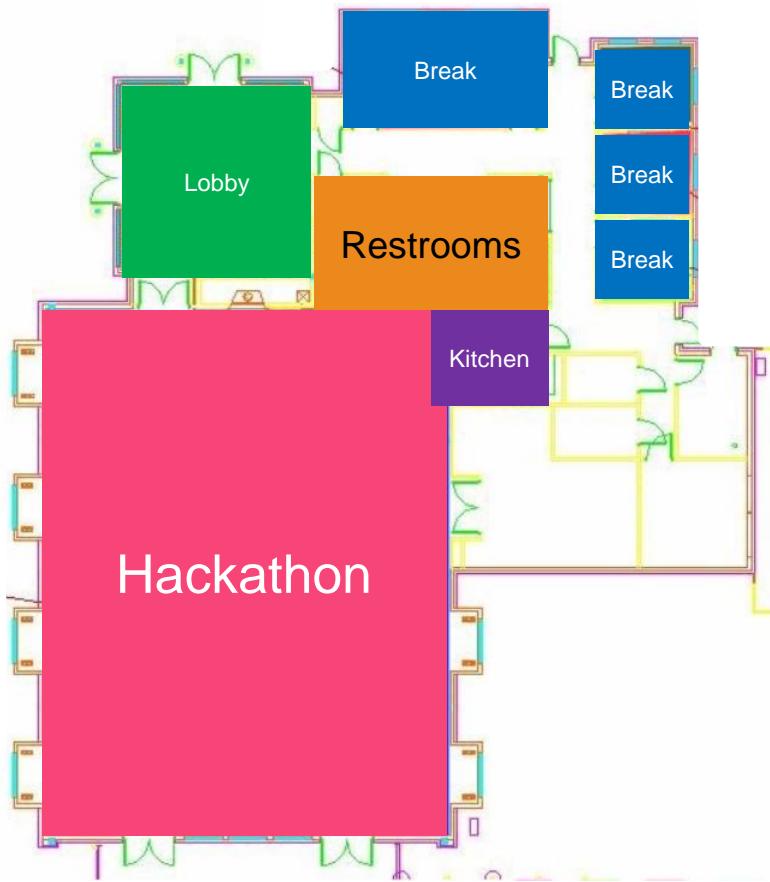


Graham Schelle

Logistics: Site Map



Logistics: Summit Retreat



Logistics: Schedule

➤ Friday:

5 PM – 7 PM	Dinner & Orientation (Silverthorne)
7 PM – 8:30 PM	PYNQ-Z1 Hand Out and Tutorials (Move to Summit Retreat)
11 PM - 12 AM	(Recommended) Project Pitch and Feedback

➤ Saturday:

12 AM – 4 AM	Mini Challenge 1
8 AM – 9 AM	Breakfast
9 AM – 1 PM	Mini Challenge 2
12 PM – 1 PM	Lunch
8 PM – 10 PM	Dinner
9:30 PM	Judging Starts
11:00 PM	Awards

Drinks and Snacks will be available throughout the event

Hackathon Prizes

Category	Value
Best Design	\$800
Runner Up	\$400
Surprise Category	\$100
Best Programmable Logic	\$100
Mini Challenge 1	Surprise Team Prize (\$150 value)
Mini Challenge 2	Surprise Team Prize (\$150 value)
Saturday Night Non-Tech Challenge	\$100

Everyone can keep their PYNQ-Z1 board!

Hackathon Judging Criteria

	<u>Score (0-10)</u>	<u>Criteria</u>	<u>Description</u>
Idea	10	Creativity	How creative or innovative is the idea behind the design?
	10	Viability	How relevant is the design's use case to a real problem?
Open Source and Reproducible	10	open-source	Are build steps provided, python packaged, and available on github as notebook and source code.
	10	reproducible	Could someone reproduce this work starting from a pynq v2.0 image and going through their notebook
Implementation	10	Feasibility	Does it actually work?
	10	Uniqueness	How well does the app leverage a variety of different open-source APIs and programmable logic

Code of Conduct

- Be Respectful
- Harassment and abuse are never tolerated
- Participants asked to stop any harassing behavior are expected to comply immediately.
- If you are being harassed, notice that someone else is being harassed, or have any other concerns, please contact Xilinx Hackathon organizers
- Full Code of Conduct in Hackathon packets

WiFi & Internet Access

➤ WiFi

- Network securebrd
- Password hackathon2017

➤ Ethernet

- No wall-plate connections
- Bridge network from laptop to PYNQ-Z1

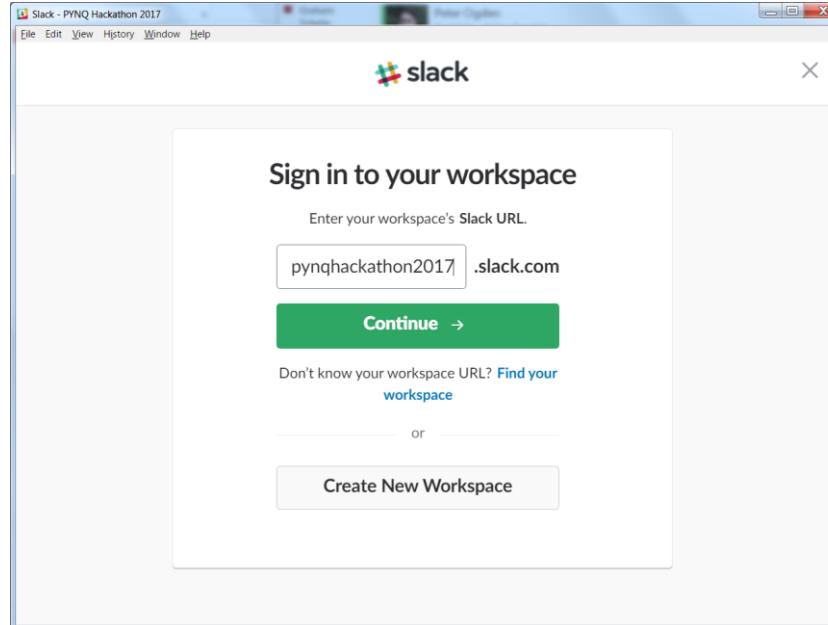
Collaboration

➤ GitHub

PYNQ Repositories

- <http://github.com/xilinx/PYNQ>
- <https://github.com/drichmond/PYNQ-Hackathon-2017>

➤ Slack

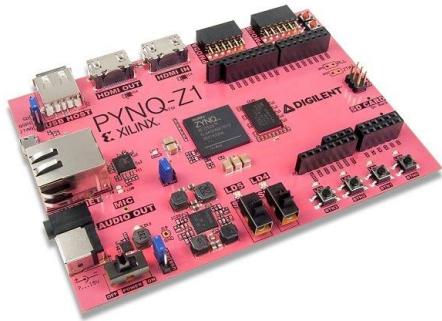


Welcome and Orientation

- Logistics
- Hackathon Theme & Goals
- Community Projects
- Head back up to Retreat!



Hackathon Themes



Build something cool

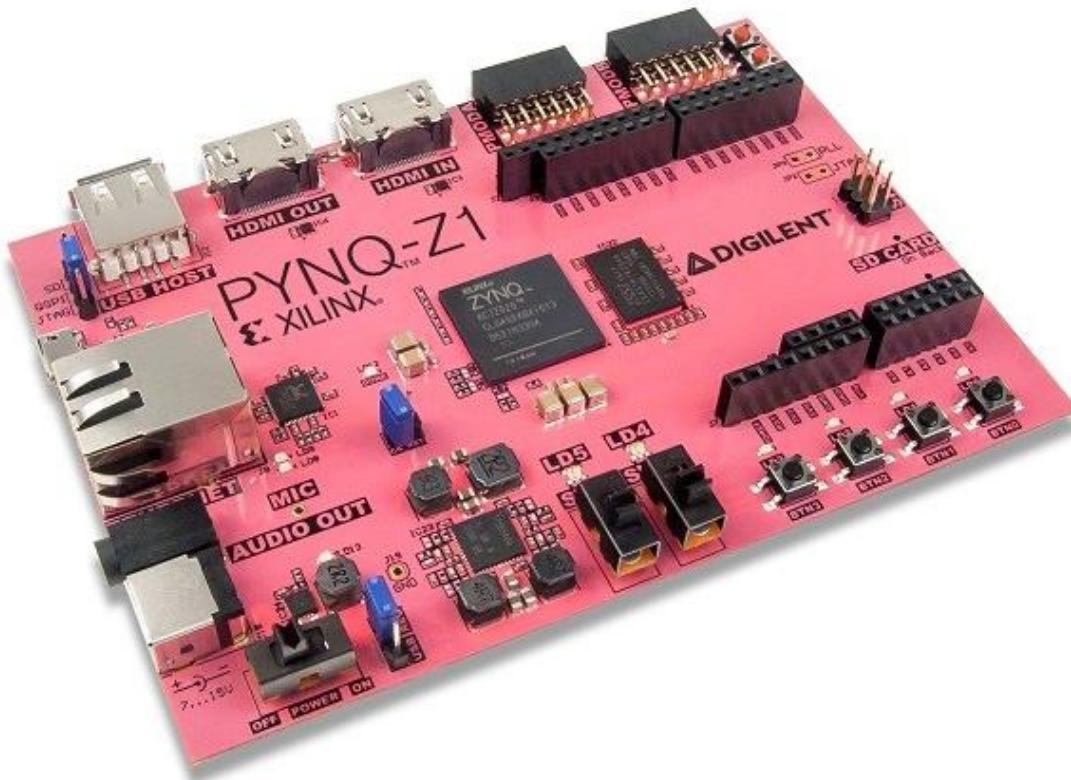


Contribute to Open Source



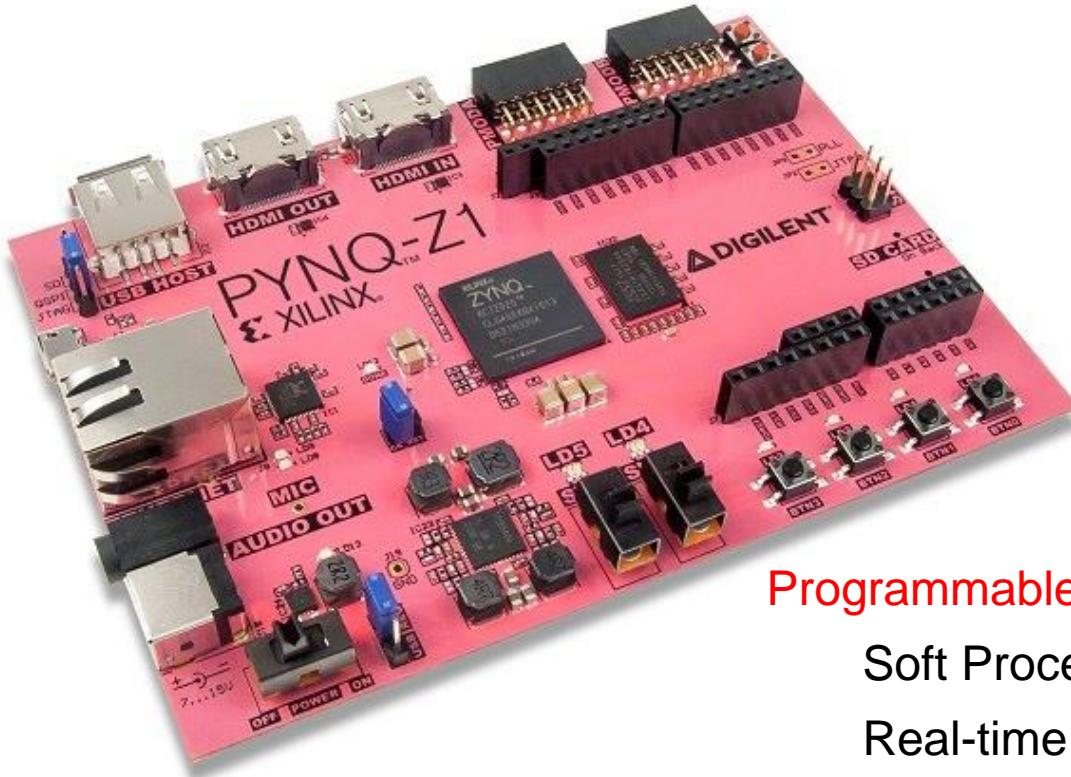
Contribute Reproducible Results

PYNQ-Z1 Board



- Zynq XC7Z020 CLG400
- 512 MB DRAM
- HDMI input, HDMI out
- Microphone input
- Audio output
- 10/100/1G Ethernet
- MicroSD slot
- USB 2.0
- 4 LEDs, 4 buttons, 2 switches
- Arduino connector
 - With additional IO
- 2 Pmod 8-pin GPIOs

PYNQ-Z1 Board



Software Preloaded

Ubuntu

Python 3.6

WiFi Connectivity

Extend it

Apt-get, pip install, git clone

Crash it

We'll burn a new SDCard

Programmable Logic Designs Preloaded

Soft Processors on Arduino & PMOD

Real-time Video Pipeline, Audio

Neural Networks

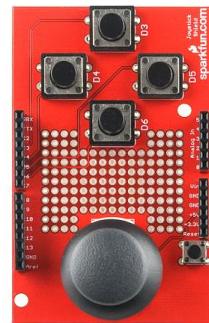
LogicTools – programmable waveforms

25+ example notebooks

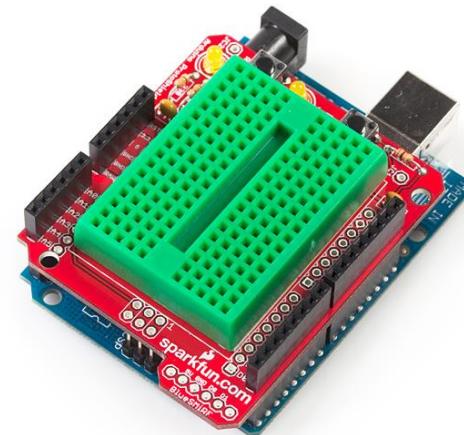
PYNQ Peripherals (1 of 4)



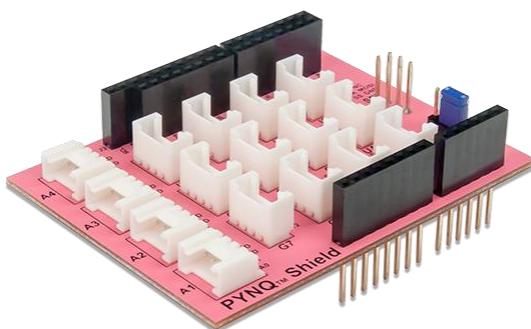
Ardumoto Shield Kit



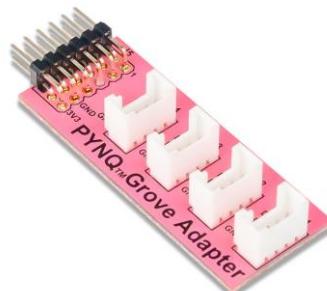
Joystick Shield Kit



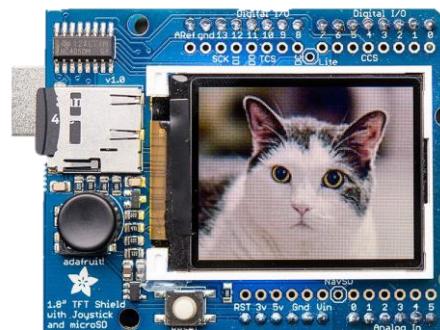
ProtoShield Kit



Arduino Grove shield



PMOD Grove Adapter



LCD18 shield

PYNQ Peripherals (2 of 4)



Light Sensor



6-Axis Accelerometer & Gyroscope



Chainable RGB LED



Ultrasonic Ranger

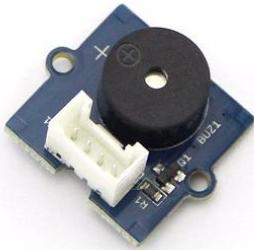


Temperature & Humidity Sensor



PIR Motion Sensor

PYNQ Peripherals (3 of 4)



Buzzer



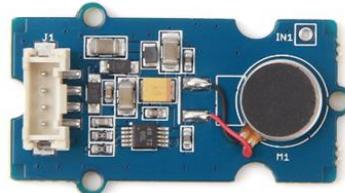
Heart Rate Monitor



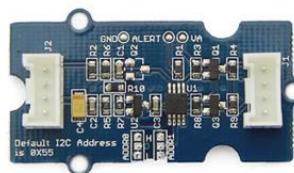
OLED Display



Color Sensor



Haptic Motor



I2C Analog-Digital
Converter



LED Bar

PYNQ Peripherals (4 of 4)



WiFi Dongle



720p USB Webcam

We'll give some project ideas when we hand out boards

PYNQ is completely open-source

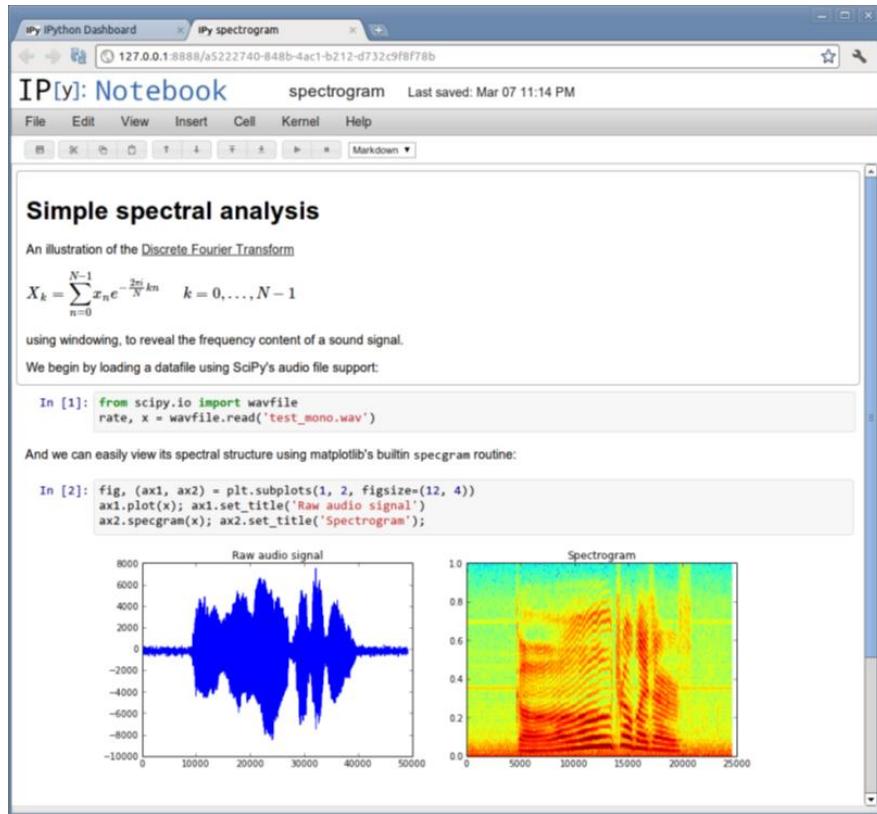
The screenshot displays three main components:

- GitHub Repository:** Shows the Xilinx/PYNQ repository on GitHub with 1,088 commits and 4 branches. It includes a list of recent pull requests and a detailed commit history.
- PYNQ Website:** The official PYNQ website (www.pynq.io) featuring a large banner image of the PYNQ-Z1 board. Below the banner, there are sections titled "What is PYNQ?" and "Who is PYNQ for?", along with a detailed description of the project's purpose and target users.
- Setup Guide:** A screenshot of the "Getting Started" section from the PYNQ-Z1 Setup Guide. It includes a video player, a "Supported Browsers" note, and a "Getting Started" section with a "Video" link and a "Setting up your PYNQ-Z1" link.

1. No Xilinx tools required to use PYNQ
2. At this hackathon, we'll strive for open-*
-source, -projects, -solutions, -collaboration



Notebooks: browser-based development ... with rich, multi-media support



- Designed for
 - Interactive, exploratory computing
 - Reproducible results
- Ideal for
 - Teaching and learning
 - Projects and research
- Rapid adoption
 - Across multiple disciplines

github.com/ipython/ipython/wiki/A-gallery-of-interesting-IPython-Notebooks

Part of a broader trend towards new, web-based IDEs



2 Caltech scientists share Nobel Prize in physics for gravitational wave discoveries



In the above spectrograms, you can see lots of excess power below ~20 Hz, as well as strong spectral lines at 500, 1000, 1500 Hz (also evident in the ASDs above). The lines at multiples of 500 Hz are the harmonics of the "violin modes" of the fibers holding up the mirrors of the LIGO interferometers.

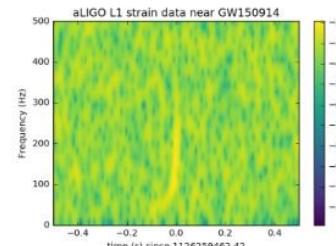
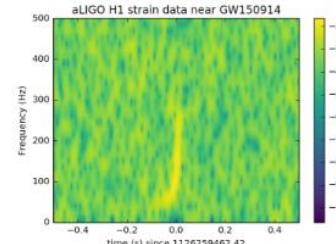
The signal is just barely visible here, at time=0 and below 500 Hz. We need to zoom in around the event time, and to the frequency range from [20, 400] Hz, and use the whitened data generated above.

```
11]: # plot the whitened data, zooming in on the signal region:
tevent = 1126259462.422           # Mon Sep 14 09:50:45 GMT 2015
deltat = 16.                      # seconds around the event
# index into the strain time series for this time interval:
indx_t = np.where((time_H1 >= tevent-deltat) & (time_H1 < tevent+deltat))

# pick a shorter FFT time interval, like 1/16 of a second:
NFFT = fs/16
# and with a lot of overlap, to resolve short-time features:
NOVL = NFFT*15/16
# and choose a window that minimizes "spectral Leakage"
# (https://en.wikipedia.org/wiki/Spectral\_leakage)
window = np.blackman(NFFT)

# Plot the H1 whitened spectrogram around the signal
plt.figure()
spec_H1, freqs, bins, im = plt.specgram(strain_H1_whiten[indx_t], NFFT=NFFT, Fs=fs, window=window,
                                         noverlap=NOVL, cmap=spec_cmap, xextent=[-deltat,deltat])
plt.xlabel('time (s) since '+str(tevent))
plt.ylabel('Frequency (Hz)')
plt.colorbar()
plt.axis([-0.5, 0.5, 0, 500])
plt.title('aLIGO H1 strain data near GW150914')
plt.savefig('GW150914_H1_spectrogram_whitened.png')

# Plot the L1 whitened spectrogram around the signal
plt.figure()
spec_L1, freqs, bins, im = plt.specgram(strain_L1_whiten[indx_t], NFFT=NFFT, Fs=fs, window=window,
                                         noverlap=NOVL, cmap=spec_cmap, xextent=[-deltat,deltat])
plt.xlabel('time (s) since '+str(tevent))
plt.ylabel('Frequency (Hz)')
plt.colorbar()
plt.axis([-0.5, 0.5, 0, 500])
plt.title('aLIGO L1 strain data near GW150914')
plt.savefig('GW150914_L1_spectrogram_whitened.png')
```



LIGO results shared as Jupyter notebooks

https://losc.ligo.org/s/events/GW150914/GW150914_tutorial.html

Welcome and Orientation

- Logistics
- Hackathon Theme & Goals
- Community Projects
- Head back up to Retreat!



Hackathons and Usergroups



Polytechnic University of Milan
100+ student hackathon



Tokyo User Group
75+ industry participants

FINN: Binary Neural Network Overlay on PYNQ

FINN: A Framework for Fast, Scalable Binarized Neural Network Inference

Yaman Umuroglu^{*†}, Nicholas J. Fraser^{*‡}, Giulio Gambardella^{*}, Michaela Blott^{*},

Philip Leong[‡], Magnus Jahre[‡], Kees Vissers^{*}

^{*}Xilinx Research Labs; [†]Norwegian University of Science and Technology; [‡]University of Sydney

ABSTRACT

Research has shown that convolutional neural networks contain significant redundancy, and high classification accuracy can be obtained even when weights and activations are reduced from floating point to binary values. In this paper, we present FINN, a framework for building fast and flexible FPGA accelerators using a flexible heterogeneous streaming architecture. By utilizing a novel set of optimizations that enable efficient mapping of binarized neural networks to hardware, we implement fully connected, convolutional and pooling layers, with per-layer compute resources being tailored to user-provided throughput requirements. On a

While the vast majority of CNNs implementations use floating point parameters, a growing body of research demonstrates this approach incorporates significant redundancy. Recently, it has been shown [7, 27, 22, 14, 32] that neural networks can classify accurately using one- or two-bit quantization for weights and activations. Such a combination of low-precision arithmetic and small memory footprint presents a unique opportunity for fast and energy-efficient image classification using Field Programmable Grid Arrays (FPGAs). FPGAs have *much* higher theoretical peak performance for binary operations compared to floating point, while the small memory footprint removes the off-chip mem-

Int. Symposium on FPGAs, Feb. 2017

- Unprecedented image classification rates
- 1,000x speed-up over Raspberry Pi3



Norwegian University of
Science and Technology



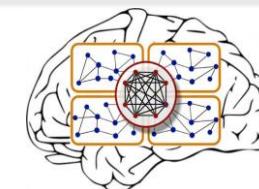
THE UNIVERSITY OF
SYDNEY



SVHN,
Road signs,
CIFAR-10



Image pre-processing
in Python



Binary Neural Network
in FPGA & ARM CPU

“cat”

Machine Learning on PYNQ cluster

Spark acceleration on FPGAs: A use case on machine learning in Pynq

Elias Koromilas, Ioannis Stamelos
Department of Electrical and Computer Engineering,
National Technical University of Athens
Athens, Greece

Christoforos Kachris
Institute of Communication and Computer Systems (ICCS/NTUA)
Athens, Greece

Dimitrios Soudris
Department of Electrical and Computer Engineering,
National Technical University of Athens
Athens, Greece

Abstract—Spark is one of the most widely used frameworks for data analytics. Spark allows fast development for several applications like machine learning, graph computations, etc. In this paper, we present Spynq: A framework for the efficient deployment of data analytics on embedded systems that are based on the heterogeneous MPSoC FPGA called Pynq. The mapping of Spark on Pynq allows that fast deployment of embedded and cyber-physical systems that are used in edge and fog computing. The proposed platform is evaluated in a typical machine learning application based on logistic regression. The performance evaluation shows that the heterogeneous FPGA-based MPSoC can achieve up to 11x speedup compared to the execution time in the ARM cores and can reduce significantly the development time of embedded and cyber-physical systems on Spark applications.

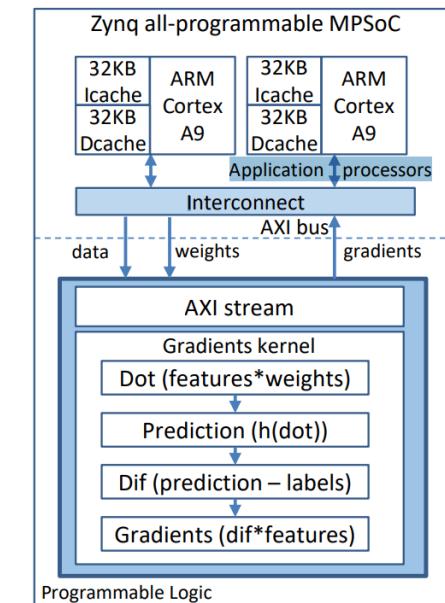
powered down in order to comply with thermal constraints [3]. One way to address this problem is through the utilization of hardware accelerators. Hardware accelerators can be used to offload the processor, increase the total throughput and reduce the energy consumption.

In this paper we present a framework for the seamlessly utilization of hardware accelerators in heterogeneous SoCs that are used to speedup the processing of Spark data analytics applications.

The main contributions of this paper are the following:

- An efficient framework for the seamlessly utilization of hardware accelerators for Spark applications

Best student paper award, MODCAST 2017



SPYNQ: Spark PYNQ cluster



National
Technical
University of
Athens



<https://www.youtube.com/watch?v=ix1NI7yq8O4>

Evaluating PYNQ for image processing

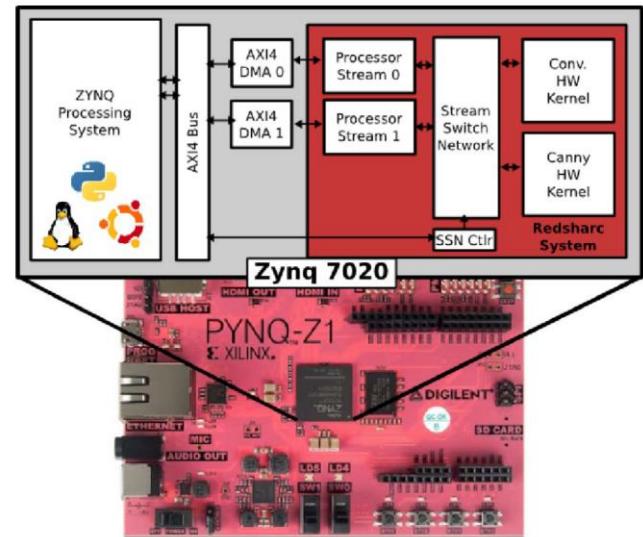
2017 IEEE 25th Annual International Symposium on Field-Programmable Custom Computing Machines

Evaluating Rapid Application Development with Python for Heterogeneous Processor-based FPGAs

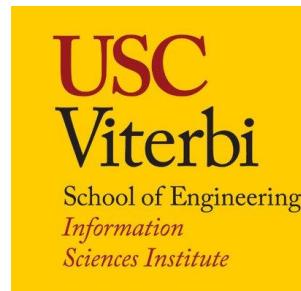
Andrew G. Schmidt, Gabriel Weisz, and Matthew French
Information Sciences Institute, University of Southern California
Email: {aschmidt, gweisz, mfrench}@isi.edu

Abstract—As modern FPGAs evolve to include more heterogeneous processing elements, such as ARM cores, it makes sense to consider these devices as processors first and FPGA accelerators second. As such, the conventional FPGA development environment must also adapt to support more software-like programming functionality. While high-level synthesis tools can help reduce FPGA development time, there still remains a large expertise gap in order to realize highly performing implementations. At a system-level the skill set necessary to integrate multiple custom IP hardware cores, interconnects, memory interfaces, and now heterogeneous processing elements is complex. Rather than drive FPGA development from the hardware up, we consider the impact of leveraging Python to accelerate application development. Python offers highly optimized

libraries and tools available to developers. Python is being used in everything from scientific computing to image processing and machine learning, and growing more each day. Making FPGAs more user friendly certainly has been an on-going effort for decades and this work does not claim to solve this problem. Instead, it looks at how entire communities have sprung up seemingly overnight around other embedded platforms, such as Raspberry Pi and Arduino. The success of these platforms stems from an inexpensive compute platform, ease of use programming environment, modularity, and a plethora of interesting and fun projects readily available to be tried, modified, and refined.



Best short paper award, FCCM 2017



“The results are highly promising , with the ability to match and exceed performances from C implementations, up to 30x speedup.”

“PYNQ is a game-changer”

Rapid Implementation of a Partially Reconfigurable Video System with PYNQ

Brad Hutchings
Dept. of Electrical and Computer Eng.
Brigham Young University
Provo, Utah 84602
Email: brad_hutchings@byu.edu

Mike Wirthlin
Dept. of Electrical and Computer Eng.
Brigham Young University
Provo, Utah 84602
Email: wirthlin@byu.edu

Abstract—Undergraduate students rapidly implement a partially-reconfigurable, real-time video processor on the Xilinx PYNQ board. The video processor performs various real-time operations including Sobel edge detection, embossing, averaging, an interactive Pong game, etc., using a separate partially-reconfigurable bit-stream for each distinct function. Selection of image-processing functions is easily accomplished via a graphical user interface that is accessed via a Jupyter notebook. As users select image-processing functions the appropriate partial bit-stream is automatically downloaded to the FPGA. The resulting system is easily and quickly developed by several undergraduate students over a period of 10 weeks with very little supervision. All code used to the project are available for download on GitHub. The productivity benefits provided by PYNQ, including Jupyter-based documentation, tutorials, and executable Python code greatly ease development effort making PYNQ an excellent FPGA platform for education.

I. INTRODUCTION

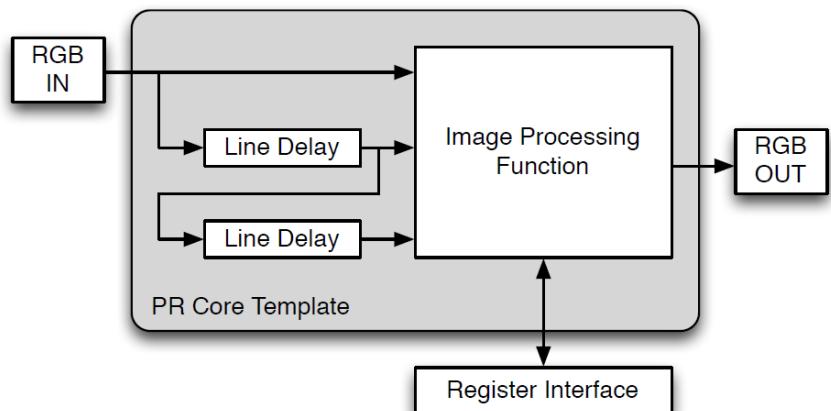
PYNQ is a new FPGA board/system designed by Xilinx and sold by Digilent that attempts to dramatically improve accessibility and productivity for FPGA-based systems. PYNQ combines a ZYNQ-7020-based FPGA board with Ubuntu Linux, extensive Jupyter[1]-based documentation with embedded Python programming examples, and a hardware “overlay” implemented in ZYNQ’s Programmable Logic (PL). The overlay

hardware functions with associated connectors, Ubuntu Linux, Jupyter notebooks, and a Python interpreter. As an FPGA board, PYNQ is unique for two reasons. First, it is immediately usable “out of the box” by users who have no FPGA experience. Second, PYNQ is completely self-documenting. Users simply connect power, insert a provided SD-card, connect an ethernet cable to their laptop or router and point a web browser

Fig. 1. Xilinx PYNQ-Z1



Partially reconfigurable video overlay

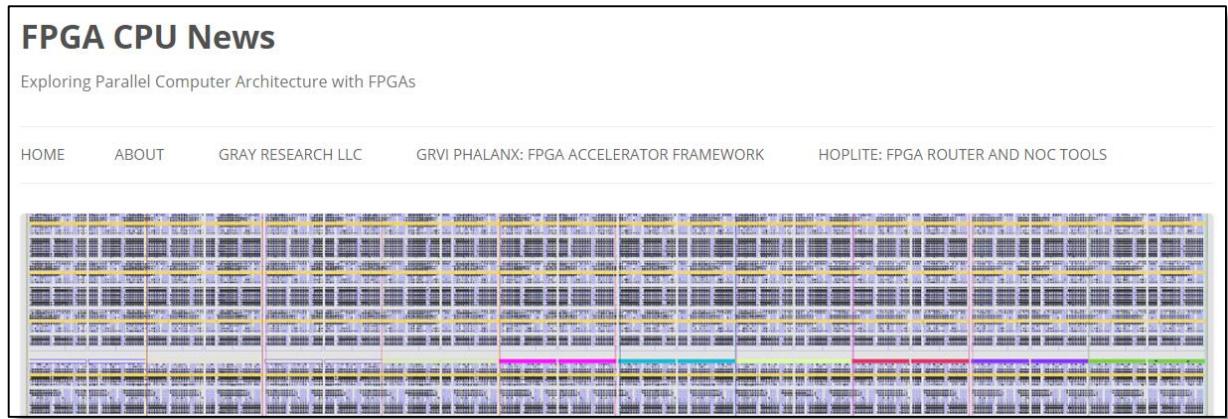


FPL2017



“After teaching students how to use FPGA-based systems for over 20 years, this is the first time we simply handed the boards to the students and **all** of them were able to use and/or add functionality to the system with no supervision.”

“A high productivity tool for experienced FPGA designers”



<http://fpga.org/2017/09/05/pynq-as-a-high-productivity-platform-for-fpga-design-and-exploration>

“Fellow FPGA designers, try Pynq. You’ll like it.
Pynq makes exploring new FPGA ideas lightweight,
fresh, fast, easy, *fun again..*”

Jan Gray ... feedback on implementing 80 x 32-bit RISC cores on PYNQ-Z1

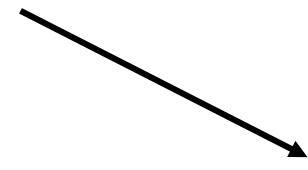
PYNQ: PYTHON PRODUCTIVITY ON ZYNQ



Home Get Started PYNQ-Z1 Bo

Community Projects

Selection of projects
and notebooks



A selection of projects from the PYNQ community is shown below. Note that some examples are built on different versions of the PYNQ image.

Binary Neural Network Xilinx Labs, NTNU, University Sydney

[BINN on Pynq](#)
The project shows how to build a binary neural network on the Zynq-7020 SoC using Python and Jupyter notebooks.



SPynq: NTUA Greece

[Py Brain Installation](#)
In this project we explore PyBrain to implement the Spynq application. PyBrain is a reinforcement learning library for Python.



Processing noisy filters PYNQ Japan user group

[はじめめる機械学習](#)
Getting Started
DSC 101
This notebook shows how to use Python to process noisy images.



Soft GPU for ZC706 Ruhr University Bochum

[Show mask & detect images](#)
From image to mask, detect, and edit
Simple example: Detect a car in a still image.
Complex example: Detect a car in a video frame, extract the car, and then detect the license plate.
Note: This is a very simple, introductory example.



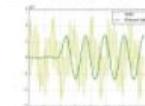
Video Processing Vectorblobx

[OpenCV Python API](#)
In this module, several filters can be applied to YUV input images. These filters are based on the OpenCV library. The code is based on the provided in Zynq-7020 Python source and Python install.



FIR filter example CU Boulder

[Implementation of a Linear Phase FIR Filter with DMA via C](#)
One of the main goals of this project is to demonstrate how to implement a linear phase FIR filter using DMA. The FIR filter is implemented using a direct form II structure. The implementation is done in C and the generated code is optimized for the Zynq-7020.



DMA and stream Tutorial

[DMA to streamed interfaces example](#)
DMA transfers from the Zynq-7020 and an AD9850 DDS board are used to demonstrate how to implement a streaming interface. One DMA will read analog samples at a rate of 200Msps. The generated DMA is then converted to a stream of digital samples. These samples are then processed by a stream interface.



CNN Example Imperial College London

[TensorFlow CNN Example](#)
TensorFlow CNN Example
This example shows how to build a TensorFlow CNN to classify handwritten digits. The code uses the MNIST dataset and the TensorFlow library.



Example Notebooks

A selection of notebook examples are shown below that are included in the PYNQ image. The notebooks contain live code, and generated output from the code can be saved in the notebook. Notebooks can be viewed as webpages, or opened on a Pynq enabled board where the code cells in a notebook can be executed.

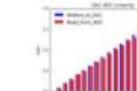
ADC waveforms

[ADC demo.ipynb](#)
An ADC demo that is designed to test the ADCs present on the Zynq-7020 SoC. It includes a graphical user interface (GUI) to control the ADCs and a Jupyter notebook to analyze the data.



DAC ADC example

[DAC ADC example.ipynb](#)
This notebook demonstrates how to interface the DAC and ADC on the Zynq-7020 SoC. It includes a GUI to control the DAC and a Jupyter notebook to analyze the ADC data.



Downloading overlays

[Downloadable Overlays](#)
This notebook demonstrates how to download an MPF overlay and execute it.

1. Installing an overlay
To generate an overlay, you need to have a ZedBoard or a Zynq-7020 SoC. You also need to have a full path to include the python package for the ZedBoard/Zynq-7020. The example of overlay generation is done with the ZedBoard. If you don't have a ZedBoard, you can use a Zynq-7020 instead.

2. Generating an overlay
If you have a ZedBoard, run the command:
`mpfgen -o ./myOverlay -c ./myOverlay.zcu106`

3. Overlaying the ZedBoard
Now you can check the download directory for the overlay:
`ls ./download`
Download to ZedBoard:
`mpfutil -d ./download`

4. Overlaying and read a single value
The code reads the MPF status file. The generated status contains an array of values.

Group Participation!

► Please introduce yourself

- Name
- School / Company
- Interests

- Do you need a team?

Welcome and Orientation

- Logistics
- Hackathon Theme & Goals
- Community Projects
- Head back up to Retreat!

